

Docker

 by Praveen Kumar

What is Virtualization?

Virtualization is a technique that allows you to create **multiple virtual environments (Virtual Machines)** on a **single physical system**.

Key points

- Uses a **hypervisor** (VMware, VirtualBox, Hyper-V).
- Each VM has its own **Operating System (OS)**.
- Multiple OSes run on the same hardware.
- Uses virtual CPU, RAM, Disk allocated from physical machine.

Problems with Virtual Machines

Although VMs are useful, they have limitations:

Heavyweight

Each VM contains a full OS → consumes more RAM, CPU, disk.

Slow Startup

Booting a full OS takes time (minutes).

Inefficient Resource Usage

Many duplicated OS kernels running on the same hardware.

App Packaging Issue

Apps inside a VM depend on the OS version and configurations.

Difficult to Move

VM images are large (GBs), hard to share.

What are Containers?

Containers are **lightweight**, **portable** environments that package an application with everything it needs—**without needing a full OS**.

Key points

- Share the **host OS kernel**.
- Start in **milliseconds**.
- Very small in size (MBs instead of GBs).
- Same behavior across machines (“It works on my machine” solved).
- Container Engines: **Docker**, **containerd**, **podman**.

Containers vs Virtual Machines

Feature	Virtual Machines (VMs)	Containers
OS	Each VM has its own full OS	Share host OS kernel
Size	Heavy (GBs)	Lightweight (MBs)
Startup Time	Slow (minutes)	Fast (seconds/milliseconds)
Resource Usage	High	Low
Portability	Difficult	Highly portable
Performance	Moderate	Almost near native

What is Docker?

Docker is a **containerization platform** that allows developers and DevOps teams to:

- Build applications as **images**
- Run applications as **containers**
- Share containers easily
- Ensure applications run the same everywhere

Docker provides:

- A container engine
- Tools to build, run, ship containers
- Docker Hub for image distribution

Why Docker Became Popular?

Lightweight

Containers are extremely small and fast.

Very Fast Deployment

Boots in seconds → perfect for microservices.

Consistency

"Works on my machine" problem solved.

Easy Packaging

Everything needed (app + dependencies) is packaged in the image.

DevOps Friendly

Perfect for CI/CD pipelines.

Scalability

Suitable for microservices and orchestration (Kubernetes).

Portability

Run the same container on: Windows, Linux, Mac, Cloud, On-premise

Docker Architecture Overview

Docker architecture follows a **client-server model** with these components:

1. **Docker Client**
2. **Docker Daemon (dockerd)**
3. **Docker Host**
4. **Docker Objects** (Images, Containers, Volumes, Networks)
5. **Docker Registry** (Docker Hub/private)

Docker Client

This is the **command-line tool** or UI used to interact with Docker.

You use:

```
docker run  
docker build  
docker pull  
docker stop
```

Client sends the request to the **Docker Daemon** via REST API.

Docker Host

The machine where Docker is installed.

It contains:

- Docker Daemon
- Containers
- Images
- Networks
- Volumes

Examples of Docker Hosts:

- Your laptop
- Cloud machine (AWS EC2)
- Linux server

Docker Daemon (dockerd)

Docker Daemon is the **brain** of Docker.

Responsibilities:

- Build images
- Run containers
- Manage networks
- Manage volumes
- Communicate with Docker Hub
- Listen for API requests from Docker Client

Acts like a background service.

Docker Registry (Docker Hub)

A **registry** stores Docker **images**.

Docker Hub:

- Public registry by Docker Inc.
- Contains official images like **nginx**, **mysql**, **node**, **python**.

Private Registries:

- AWS ECR
- GitHub Container Registry
- Azure ACR
- Google GCR

Docker Images

A **Docker Image** is a **read-only blueprint** used to create containers.

Key characteristics

- Contains:
 - Application code
 - Runtime (Python/Node/Java)
 - Libraries
 - System tools
- Built using a **Dockerfile**
- Images are made up of **layers**
 - Each instruction (FROM, COPY, RUN) creates a new layer
- Images are immutable (cannot be changed once created)

Docker Containers

A **Docker Container** is a **running instance** of a Docker Image.

Key characteristics

- Lightweight and fast
- Includes:
 - App binaries
 - Dependencies
- Uses **copy-on-write** for filesystem changes
- Has its own:
 - Process space
 - Network interface
 - File system

Docker Registry, Repository, Tags

Docker Registry

A registry is a **storage system for Docker images**.

Examples:

- Docker Hub (public)
- AWS ECR
- GitHub Container Registry
- Azure ACR

Docker Repository

A **repository** is a **collection of images for the same application** with different versions.

Basic Docker Information Commands

Command	Description
<code>docker --version</code>	Shows installed Docker version
<code>docker version</code>	Shows Client & Daemon version details
<code>docker info</code>	Displays Docker system information (containers, images, storage, etc.)

Docker Image Commands

Command	Description
docker images	List all images available locally
docker pull <image>	Download image from Docker Hub/Registry
docker build -t <name> .	Build image using Dockerfile
docker rmi <image>	Remove an image
docker image prune	Remove unused images
docker inspect <image>	Shows detailed metadata of an image
docker tag <src> <target>	Tag an image with a new name/version

Docker Container Commands

Command	Description
<code>docker run <image></code>	Creates + starts a container
<code>docker run -d <image></code>	Run container in detached (background) mode
<code>docker ps</code>	List running containers
<code>docker ps -a</code>	List all containers (including stopped ones)
<code>docker start <container></code>	Start a stopped container
<code>docker stop <container></code>	Stop a running container
<code>docker restart <container></code>	Restart container

```
docker kill <container>
```

Force stop the container

```
docker rm <container>
```

Remove a stopped container

```
docker logs <container>
```

Show container logs

```
docker exec -it <container> bash
```

Execute a command inside running container

```
docker inspect <container>
```

Detailed container information

```
docker stats
```

Real-time resource usage of containers

Docker Volume Commands

Command	Description
<code>docker volume create <name></code>	Create a new volume
<code>docker volume ls</code>	List volumes
<code>docker volume inspect <name></code>	Show details of a volume
<code>docker volume rm <name></code>	Remove a volume
<code>docker volume prune</code>	Remove all unused volumes

Docker Network Commands

Command	Description
<code>docker network ls</code>	List available Docker networks
<code>docker network create <name></code>	Create a custom network
<code>docker network inspect <name></code>	Show network details
<code>docker network rm <name></code>	Remove a network
<code>docker run --network=<name></code>	Run container in a specific network

What is a Dockerfile?

A **Dockerfile** is a simple text file containing a **set of instructions** used to build a Docker Image.

It automates the steps:

- install software
- copy source code
- set environment variables
- run commands
- expose ports
- define how the container starts

In simple words:

Dockerfile = Recipe

Image = Cake

Container = A piece of cake you eat (running instance)

Why Dockerfile?

- Reproducible builds
- Portable across environments
- Easy to version control
- Removes manual steps
- Standard way for CI/CD

Basic Structure of a Dockerfile

A Dockerfile is a list of instructions executed **from top to bottom**.

Example:

```
FROM python:3.10
WORKDIR /app
COPY ..
RUN pip install -r requirements.txt
EXPOSE 5000
CMD ["python", "app.py"]
```

Dockerfile Instructions

FROM - Every image starts from a base image (like python, ubuntu, node, nginx).

WORKDIR - Sets the working directory inside the container.

COPY - Copies files from your machine → to container.

ADD - Similar to COPY but supports:

- URL downloads
- Automatic extraction of tar files

RUN - Runs a command **during image build**.

CMD - Defines the command that runs when the container starts.

ENTRYPOINT - Similar to CMD but **cannot be overridden easily**.

EXPOSE - Tells Docker which port the container will use.

ENV - Sets environment variables inside the container.

VOLUME - For persistent data storage.

Building an Image Using Dockerfile

Step 1: Create Dockerfile

Step 2: Build the image:

```
docker build -t myapp .
```

Step 3: Run a container:

```
docker run -p 5000:5000 myapp
```

Simple Python App (Flask)

Dockerfile:

```
FROM python:3.10
WORKDIR /app
COPY ..
RUN pip install flask
EXPOSE 5000
CMD ["python", "app.py"]
```

Node.js App

Dockerfile

```
FROM node:18
WORKDIR /usr/src/app
COPY package*.json .
RUN npm install
COPY ..
EXPOSE 3000
CMD ["npm", "start"]
```

Best Practices

- ✓ Use small images (like alpine)
- ✓ Keep the number of layers small
- ✓ Use .dockerignore
- ✓ Pin version numbers (avoid latest)
- ✓ Use multi-stage builds for production
- ✓ Combine RUN commands
- ✓ Do not install unnecessary packages

What Are Docker Volumes?

A **Docker Volume** is a special storage mechanism used by Docker to **persist data** generated and used by containers.

Because containers are **temporary**, once a container is removed:

- All data stored inside the container is **lost**.

To prevent data loss, Docker provides **Volumes**.

✓ Summary (easy definition)

Volumes store data outside the container's filesystem so the data will not be deleted when the container is removed.

Why Do We Need Volumes?

Containers are:

- *Ephemeral* (temporary)
- *Stateless* by default

If you store data inside a running container (like in `/var/lib/mysql` for MySQL), deleting or recreating the container deletes all the data.

Example:

- Run MySQL container
- Store 10 rows
- Delete container
- All data gone 

To avoid this → **use a volume** 

Benefits of Docker Volumes

Benefit	Explanation
Data Persistence	Data survives even after container removal
Sharing Data	Multiple containers can access the same volume
Better Performance	Faster I/O than bind mounts
Backup & Restore	Easy to back up volume directories
Decoupling Data from Containers	Safe upgrades without losing data

Types of Docker Storage

Docker provides 3 types of storage:

1 Volumes (Recommended)

- Managed by Docker
- Stored under `/var/lib/docker/volumes/`
- Best for production
- Works on Linux, Windows, Mac equally

2 Bind Mounts

- Maps a folder from the host machine → into container
- Used for local development
- Example: Mapping source code from your laptop

```
docker run -v C:/code:/app node
```

3 tmpfs Mounts

- Data stored in RAM only
- Fastest but temporary
- Good for sensitive data

Types of Docker Volumes

A) Anonymous Volumes

Created automatically when you use `-v /path`.

Not easy to reference again.

Example:

```
docker run -v /app/data nginx
```

B) Named Volumes (most important)

Explicitly created by the user.

Example:

```
docker volume create mydata  
docker run -v mydata:/var/lib/mysql mysql
```

Where Are Volumes Stored?

On Linux:

```
/var/lib/docker/volumes/<volume_name>/_data/
```

Docker Volume Commands

1. Create a volume docker volume create myvolume
2. List volumes docker volume ls
3. Inspect a volume docker volume inspect myvolume
4. Remove volume docker volume rm myvolume
5. Remove unused volumes docker volume prune

Named Volume Exam ple

Step 1: Create a volume

```
docker volume create mydata
```

Step 2: Run a container using the volume

```
docker run -d \  
  --name mycontainer \  
  -v mydata:/app/data \  
  nginx
```

This mounts the volume **mydata** inside the container at **/app/data**.

Step 3: Copy some data into the container

Step 3: Copy some data into the container

```
docker exec -it mycontainer bash  
cd /app/data  
echo "Hello Volume" > file.txt  
exit
```

Step 4: Stop & remove the container

```
docker rm -f mycontainer
```

Step 5: Run another container with same volume

```
docker run -it --name newcontainer -v mydata:/app/data ubuntu bash
```

Inside this new container:

```
cat /app/data/file.txt
```

You will see your data!

Volumes persist even after container deletion.

Anonymous Volume Example

This creates a volume **without a name**:

Run container

```
docker run -d --name testcontainer -v /app/data nginx
```

Docker automatically creates a volume with a random name (e.g., 73f8d2aca91a34...)

Dockerfile Example Using Volumes

Create a file named **Dockerfile**:

```
FROM ubuntu:latest
RUN mkdir /appdata
VOLUME ["/appdata"]
CMD ["bash"]
```

Build the image:

```
docker build -t volume-demo .
```

Run the container:

```
docker run -it --name vol1 volume-demo
```

Inside container:

```
docker exec -it <containername> /bin/bash  
cd /appdata  
echo "Hello Dockerfile Volume" > test.txt
```

Exit container and remove it:

```
docker rm -f vol1
```

Run another container:

```
docker run -it --name vol2 volume-demo
```

✓ Your file test.txt is still present, because Docker automatically creates an anonymous volume for /appdata.

What is Docker Compose?

Docker Compose is a tool that allows you to **define and run multi-container Docker applications** using a single configuration file called:

`docker-compose.yml`

With Compose, you can:

- Start multiple containers together
- Configure networks & volumes automatically
- Link services together using service names
- Scale services
- Stop and remove everything with one command

Why Docker Compose?

Without Docker Compose, if you want to run a multi-container app (e.g., web + database + cache), you would manually run commands like:

```
docker run -d -p 8080:80 webapp  
docker run -d mysql  
docker run -d redis  
...
```

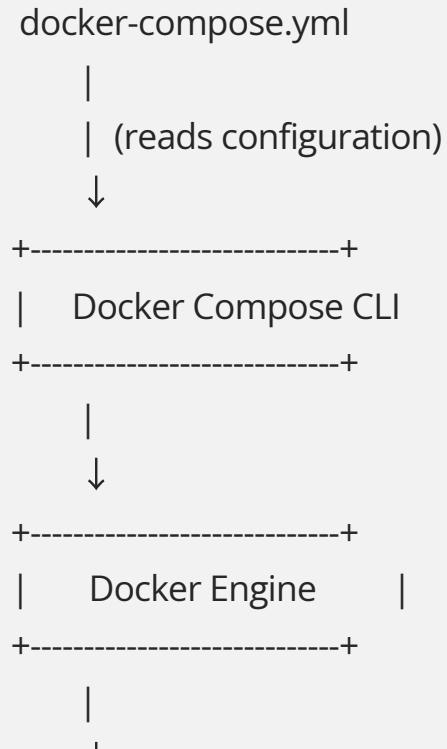
This becomes **complicated** and **unmanageable** as the number of containers grows.

With Docker Compose:

```
docker compose up
```

Everything starts automatically based on the compose file.

Docker Compose Architecture (Simple Diagram)



Services → Networks → Volumes → Containers

docker-compose.yml Structure Overview

A Compose file has the following sections:

```
version: "3.9"
```

```
services:
```

```
  serviceA:
```

```
    image: ...
```

```
    ports:
```

```
    volumes:
```

```
    environment:
```

```
    depends_on:
```

```
    networks:
```

```
networks:
```

```
  mynetwork:
```

```
volumes:
```

```
  myvolume:
```

omplete Example 1— Basic Web + Database (Very Simple)

File: docker-compose.yml

```
version: "3.9"

services:
  web:
    image: nginx
    ports:
      - "8080:80"

  db:
    image: mysql:5.7
    environment:
      MYSQL_ROOT_PASSWORD: root123
```

Run app:

```
docker compose up
```

Run in background:

```
docker compose up -d
```

Stop app:

```
docker compose down
```

Complete Example 2 — Web + MySQL + Volume + Network

This is the **most commonly used structure**.

```
version: "3.9"

services:

backend:
  build: ./backend
  container_name: backend-app
  ports:
    - "5000:5000"
  depends_on:
    - mysql
  environment:
    DB_HOST: mysql
    DB_USER: root
    DB_PASS: root123
  networks:
    - mynet
```

```
mysql:  
  image: mysql:8  
  container_name: mysql-db  
  environment:  
    MYSQL_ROOT_PASSWORD: root123  
    MYSQL_DATABASE: testdb  
  volumes:  
    - mysql_data:/var/lib/mysql  
  networks:  
    - mynet
```

```
volumes:  
  mysql_data:  
  
networks:  
  mynet:
```

Useful Docker Compose Commands

Command	Description
docker compose up	Start all services
docker compose up -d	Start in background
docker compose down	Stop + remove containers + network
docker compose down -v	Remove volumes also
docker compose ps	List running services
docker compose logs	View logs of all services
docker compose logs backend	Logs for a single service

```
docker compose stop
```

Stop but do not remove

```
docker compose start
```

Start previously stopped services

```
docker compose build
```

Rebuild images

```
docker compose exec backend bash
```

Enter a container

```
docker compose restart
```

Restart services

Life Cycle of Docker Compose (Flow)

docker-compose.yml



docker compose up



Creates networks → Creates volumes → Starts containers → Connects them

