

# Docker

 by Praveen Kumar

# What is Virtualization?

Virtualization is a technique that allows you to create **multiple virtual environments (Virtual Machines)** on a **single physical system**.

## Key points

- Uses a **hypervisor** (VMware, VirtualBox, Hyper-V).
- Each VM has its own **Operating System (OS)**.
- Multiple OSes run on the same hardware.
- Uses virtual CPU, RAM, Disk allocated from physical machine.

# Problems with Virtual Machines

Although VMs are useful, they have limitations:

## **Heavyweight**

Each VM contains a full OS → consumes more RAM, CPU, disk.

## **Slow Startup**

Booting a full OS takes time (minutes).

## **Inefficient Resource Usage**

Many duplicated OS kernels running on the same hardware.

## **App Packaging Issue**

Apps inside a VM depend on the OS version and configurations.

## **Difficult to Move**

VM images are large (GBs), hard to share.

# What are Containers?

Containers are **lightweight**, **portable** environments that package an application with everything it needs—**without needing a full OS**.

## Key points

- Share the **host OS kernel**.
- Start in **milliseconds**.
- Very small in size (MBs instead of GBs).
- Same behavior across machines (“It works on my machine” solved).
- Container Engines: **Docker**, **containerd**, **podman**.

# Containers vs Virtual Machines

| Feature        | Virtual Machines (VMs)             | Containers                  |
|----------------|------------------------------------|-----------------------------|
| OS             | Each VM has its <b>own full OS</b> | Share <b>host OS kernel</b> |
| Size           | Heavy (GBs)                        | Lightweight (MBs)           |
| Startup Time   | Slow (minutes)                     | Fast (seconds/milliseconds) |
| Resource Usage | High                               | Low                         |
| Portability    | Difficult                          | Highly portable             |
| Performance    | Moderate                           | Almost near native          |

# What is Docker?

Docker is a **containerization platform** that allows developers and DevOps teams to:

- Build applications as **images**
- Run applications as **containers**
- Share containers easily
- Ensure applications run the same everywhere

## Docker provides:

- A container engine
- Tools to build, run, ship containers
- Docker Hub for image distribution

# Why Docker Became Popular?

## **Lightweight**

Containers are extremely small and fast.

## **Very Fast Deployment**

Boots in seconds → perfect for microservices.

## **Consistency**

"Works on my machine" problem solved.

## **Easy Packaging**

Everything needed (app + dependencies) is packaged in the image.

## **DevOps Friendly**

**Perfect for CI/CD pipelines.**

## **Scalability**

Suitable for microservices and orchestration (Kubernetes).

## **Portability**

Run the same container on: Windows, Linux, Mac, Cloud, On-premise

# Docker Architecture Overview

Docker architecture follows a **client-server model** with these components:

1. **Docker Client**
2. **Docker Daemon (dockerd)**
3. **Docker Host**
4. **Docker Objects** (Images, Containers, Volumes, Networks)
5. **Docker Registry** (Docker Hub/private)

Docker Client → Docker Daemon → Host OS → Containers



Docker Registry (Docker Hub)

# Docker Client

This is the **command-line tool** or UI used to interact with Docker.

## You use:

```
docker run  
docker build  
docker pull  
docker stop
```

Client sends the request to the **Docker Daemon** via REST API.

# Docker Host

The machine where Docker is installed.

It contains:

- Docker Daemon
- Containers
- Images
- Networks
- Volumes

## Examples of Docker Hosts:

- Your laptop
- Cloud machine (AWS EC2)
- Linux server

# Docker Daemon (dockerd)

Docker Daemon is the **brain** of Docker.

## Responsibilities:

- Build images
- Run containers
- Manage networks
- Manage volumes
- Communicate with Docker Hub
- Listen for API requests from Docker Client

Acts like a background service.

# Docker Registry (Docker Hub)

A **registry** stores Docker **images**.

## Docker Hub:

- Public registry by Docker Inc.
- Contains official images like **nginx**, **mysql**, **node**, **python**.

## Private Registries:

- AWS ECR
- GitHub Container Registry
- Azure ACR
- Google GCR

# Docker Images

A **Docker Image** is a **read-only blueprint** used to create containers.

## Key characteristics

- Contains:
  - Application code
  - Runtime (Python/Node/Java)
  - Libraries
  - System tools
- Built using a **Dockerfile**
- Images are made up of **layers**
  - Each instruction (FROM, COPY, RUN) creates a new layer
- Images are immutable (cannot be changed once created)

# Docker Containers

A **Docker Container** is a **running instance** of a Docker Image.

## Key characteristics

- Lightweight and fast
- Includes:
  - App binaries
  - Dependencies
- Uses **copy-on-write** for filesystem changes
- Has its own:
  - Process space
  - Network interface
  - File system

# Docker Registry, Repository, Tags

## Docker Registry

A registry is a **storage system for Docker images**.

Examples:

- Docker Hub (public)
- AWS ECR
- GitHub Container Registry
- Azure ACR

## Docker Repository

A **repository** is a **collection of images for the same application** with different versions.

# Basic Docker Information Commands

| Command                       | Description                                                            |
|-------------------------------|------------------------------------------------------------------------|
| <code>docker --version</code> | Shows installed Docker version                                         |
| <code>docker version</code>   | Shows Client & Daemon version details                                  |
| <code>docker info</code>      | Displays Docker system information (containers, images, storage, etc.) |

# Docker Image Commands

| Command                   | Description                             |
|---------------------------|-----------------------------------------|
| docker images             | List all images available locally       |
| docker pull <image>       | Download image from Docker Hub/Registry |
| docker build -t <name> .  | Build image using Dockerfile            |
| docker rmi <image>        | Remove an image                         |
| docker image prune        | Remove unused images                    |
| docker inspect <image>    | Shows detailed metadata of an image     |
| docker tag <src> <target> | Tag an image with a new name/version    |

# Docker Container Commands

| Command                                       | Description                                  |
|-----------------------------------------------|----------------------------------------------|
| <code>docker run &lt;image&gt;</code>         | Creates + starts a container                 |
| <code>docker run -d &lt;image&gt;</code>      | Run container in detached (background) mode  |
| <code>docker ps</code>                        | List running containers                      |
| <code>docker ps -a</code>                     | List all containers (including stopped ones) |
| <code>docker start &lt;container&gt;</code>   | Start a stopped container                    |
| <code>docker stop &lt;container&gt;</code>    | Stop a running container                     |
| <code>docker restart &lt;container&gt;</code> | Restart container                            |

```
docker kill <container>
```

Force stop the container

```
docker rm <container>
```

Remove a stopped container

```
docker logs <container>
```

Show container logs

```
docker exec -it <container> bash
```

Execute a command inside running container

```
docker inspect <container>
```

Detailed container information

```
docker stats
```

Real-time resource usage of containers

# Docker Volume Commands

| Command                                         | Description               |
|-------------------------------------------------|---------------------------|
| <code>docker volume create &lt;name&gt;</code>  | Create a new volume       |
| <code>docker volume ls</code>                   | List volumes              |
| <code>docker volume inspect &lt;name&gt;</code> | Show details of a volume  |
| <code>docker volume rm &lt;name&gt;</code>      | Remove a volume           |
| <code>docker volume prune</code>                | Remove all unused volumes |

# Docker Network Commands

| Command                                          | Description                         |
|--------------------------------------------------|-------------------------------------|
| <code>docker network ls</code>                   | List available Docker networks      |
| <code>docker network create &lt;name&gt;</code>  | Create a custom network             |
| <code>docker network inspect &lt;name&gt;</code> | Show network details                |
| <code>docker network rm &lt;name&gt;</code>      | Remove a network                    |
| <code>docker run --network=&lt;name&gt;</code>   | Run container in a specific network |

# What is a Dockerfile?

A **Dockerfile** is a simple text file containing a **set of instructions** used to build a Docker Image.

It automates the steps:

- install software
- copy source code
- set environment variables
- run commands
- expose ports
- define how the container starts

**In simple words:**

**Dockerfile = Recipe**

**Image = Cake**

**Container = A piece of cake you eat (running instance)**

# Why Dockerfile?

- Reproducible builds
- Portable across environments
- Easy to version control
- Removes manual steps
- Standard way for CI/CD

# Basic Structure of a Dockerfile

A Dockerfile is a list of instructions executed **from top to bottom**.

Example:

```
FROM python:3.10
WORKDIR /app
COPY ..
RUN pip install -r requirements.txt
EXPOSE 5000
CMD ["python", "app.py"]
```

# Dockerfile Instructions

FROM - Every image starts from a base image (like python, ubuntu, node, nginx).

WORKDIR - Sets the working directory inside the container.

COPY - Copies files from your machine → to container.

ADD - Similar to COPY but supports:

- URL downloads
- Automatic extraction of tar files

RUN - Runs a command **during image build**.

CMD - Defines the command that runs when the container starts.

ENTRYPOINT - Similar to CMD but **cannot be overridden easily**.

EXPOSE - Tells Docker which port the container will use.

**ENV** - Sets environment variables inside the container.

**VOLUME** - For persistent data storage.

### **Building an Image Using Dockerfile**

## **Step 1: Create Dockerfile**

Step 2: Build the image:

```
docker build -t myapp .
```

Step 3: Run a container:

```
docker run -p 5000:5000 myapp
```

# Simple Python App (Flask)

## Dockerfile:

```
FROM python:3.10
WORKDIR /app
COPY ..
RUN pip install flask
EXPOSE 5000
CMD ["python", "app.py"]
```

# Node.js App

## Dockerfile

```
FROM node:18
WORKDIR /usr/src/app
COPY package*.json .
RUN npm install
COPY ..
EXPOSE 3000
CMD ["npm", "start"]
```

# Best Practices

- ✓ Use small images (like alpine)
- ✓ Keep the number of layers small
- ✓ Use .dockerignore
- ✓ Pin version numbers (avoid latest)
- ✓ Use multi-stage builds for production
- ✓ Combine RUN commands
- ✓ Do not install unnecessary packages

# What Are Docker Volumes?

A **Docker Volume** is a special storage mechanism used by Docker to **persist data** generated and used by containers.

Because containers are **temporary**, once a container is removed:

- All data stored inside the container is **lost**.

To prevent data loss, Docker provides **Volumes**.

## ✓ Summary (easy definition)

Volumes store data outside the container's filesystem so the data will not be deleted when the container is removed.

# Why Do We Need Volumes?

Containers are:

- *Ephemeral* (temporary)
- *Stateless* by default

If you store data inside a running container (like in `/var/lib/mysql` for MySQL), deleting or recreating the container deletes all the data.

## Example:

- Run MySQL container
- Store 10 rows
- Delete container
- All data gone 

To avoid this → **use a volume** 

# Benefits of Docker Volumes

| Benefit                                | Explanation                                    |
|----------------------------------------|------------------------------------------------|
| <b>Data Persistence</b>                | Data survives even after container removal     |
| <b>Sharing Data</b>                    | Multiple containers can access the same volume |
| <b>Better Performance</b>              | Faster I/O than bind mounts                    |
| <b>Backup &amp; Restore</b>            | Easy to back up volume directories             |
| <b>Decoupling Data from Containers</b> | Safe upgrades without losing data              |

# Types of Docker Storage

Docker provides 3 types of storage:

## 1 Volumes (Recommended)

- Managed by Docker
- Stored under `/var/lib/docker/volumes/`
- Best for production
- Works on Linux, Windows, Mac equally

## 2 Bind Mounts

- Maps a folder from the host machine → into container
- Used for local development
- Example: Mapping source code from your laptop

```
docker run -v C:/code:/app node
```

## 3 tmpfs Mounts

- Data stored in RAM only
- Fastest but temporary
- Good for sensitive data

# Types of Docker Volumes

## A) Anonymous Volumes

Created automatically when you use `-v /path`.

Not easy to reference again.

Example:

```
docker run -v /app/data nginx
```

## B) Named Volumes (most important)

Explicitly created by the user.

Example:

```
docker volume create mydata  
docker run -v mydata:/var/lib/mysql mysql
```

# Where Are Volumes Stored?

On Linux:

```
/var/lib/docker/volumes/<volume_name>/_data/
```

# Docker Volume Commands

1. Create a volume docker volume create myvolume
2. List volumes docker volume ls
3. Inspect a volume docker volume inspect myvolume
4. Remove volume docker volume rm myvolume
5. Remove unused volumes docker volume prune